


# Basics of Networking 2: DNS

Daniel STAN



May 27, 2024

Original content by Eric Milard  *Éléments de réseau 2* Contrib and Speakers: Daniel Stan, Nidà Meddouri, Elloh Adja, Suzana Dedefa, Christian Diaconu

Version: v1.0, feedback and remarks to: [daniel.stan@epita.fr](mailto:daniel.stan@epita.fr)

# Outline

- ① Domain Names at a glance
- ② Iterative Queries on a system
- ③ Recursive Resolution

# Motivation: When I connect to a server, what happens?

```
dstan@flan:~$ ping epita.fr
PING epita.fr (2606:4700:20::ac43:4796) 56 data bytes
64 bytes from 2606:4700:20::ac43:4796 (2606:4700:20::ac43:4796): icmp_seq=1 ttl=51 time=12.5 ms
64 bytes from 2606:4700:20::ac43:4796 (2606:4700:20::ac43:4796): icmp_seq=2 ttl=51 time=11.4 ms
64 bytes from 2606:4700:20::ac43:4796 (2606:4700:20::ac43:4796): icmp_seq=3 ttl=51 time=11.2 ms
64 bytes from 2606:4700:20::ac43:4796 (2606:4700:20::ac43:4796): icmp_seq=4 ttl=51 time=10.7 ms
64 bytes from 2606:4700:20::ac43:4796 (2606:4700:20::ac43:4796): icmp_seq=5 ttl=51 time=10.3 ms
```

- Traduction mechanism from a **name** to an **IP address**

- ▶ ...and from IP back to name (reverse)
- ▶ and to store arbitrary data
- ▶ This is called a **resolution**, or **lookup**.

- Why:

- ▶ Easier to **remember** a name than an IP.
- ▶ Flexibility: since an IP address may change, do not **hardcode** it. Ex: **DynDNS**.
- ▶ Resilience: store more than one IP for a name. Ex: **Round Robin**, **IPv4/IPv6 dual stack**

# The entry hierarchy



Root

TLD  
(top Level Domain)



generic



geographic

Type or category of the domain

SLD  
(2<sup>nd</sup> Level Domain)



Organisation Name

Subdomain



Subcategory, expected services

# Fully Qualified Domain Name (FQDN)

A FQDN is **globally unique**.

- It contains at least a TLD and a SLD : *epita.fr*, *ionis-group.com*...
- It can contain **arbitrary** subdirectories: *cri.epita.fr*, *pie.cri.epita.fr*, *fleet.pie.crie.epita.fr*
- If a domain name is not a FQDN, it may be completed by the system using some default suffix.

# In practice: How to buy a domain?

The screenshot shows the Gandi.net website with the search bar containing 'net2-is-fun.com'. Below the search bar, there are search filters and a list of results. The first result is 'net2-is-fun.com' with a price of 17,95 € for 1 year, plus 23,99 € per year. Other results include 'net2-is-fun.fr', 'net2-is-fun.org', and 'net2-is.fun'.

**gandi.net** DOMAINES ▾ HÉBERGEMENT ▾ CERTIFICATS SSL

net2-is-fun.com

INCLUS AVEC CHAQUE NOM DE DOMAINE

- ✓ Alias et redirections illimités
- ✓ 1 certificat SSL Standard pour un an
- ✓ Serveurs de nom LiveDNS (Anycast + DNSSEC)

894 résultats

Domain	Price	Yearly Price	Action
net2-is-fun.com	17,95 €	puis 23,99 € /an	
net2-is-fun.fr	14,40 €	puis 17,99 € /an	
net2-is-fun.org	10,79 €	puis 23,99 € /an	
Pack .com + .fr + .org	50,50 €	35,23 €	Ajouter au panier
net-is.fun	3,60 €	puis 34,66 € /an	
net2-is.fun	3,60 €	puis 34,66 € /an	

Usually, one buys an **SLD**: *epita.fr*, *ionis-group.com*, *google.com* to an ICANN accredited *registrar* ( *reg-istraire*).

← example of a Registrar interface.  
Usually, if you own a domain name, you can register **any subdomain**.

Public database of domain name owners: **whois**.

Exercice: who owns *epita.fr*? *tudo.re*? *gouv.fr*?

## Iterative Queries

# A software for DNS queries

DNS Protocol: 53/udp (but 53/tcp too)

Plenty of clients: host, nslookup, dig

```
dstan@flan:~$ host epita.fr 8.8.8.8
Using domain server:
Name: 8.8.8.8
Address: 8.8.8.8#53
Aliases:

epita.fr has address 104.26.7.225
...

dstan@flan:~$ nslookup epita.fr 8.8.8.8
...

dstan@flan:~$ dig epita.fr @8.8.8.8

; <<>> DiG 9.18.18-0ubuntu0.22.04.2-Ubuntu <<>> epita.fr @8.8.8.8
...
;; ANSWER SECTION:
epita.fr.      300 IN  A 104.26.6.225
```





# What server to use? Resolvconf file

Any system stores<sup>1</sup> at least one known and reliable DNS server: `/etc/resolv.conf`

```
dstan@flan:~$ cat /etc/resolv.conf
# Dynamic resolv.conf(5) file for glibc resolver(3) generated by resolvconf(8)
#     DO NOT EDIT THIS FILE BY HAND — YOUR CHANGES WILL BE OVERWRITTEN
# 127.0.0.53 is the systemd-resolved stub resolver.
# run "systemd-resolve --status" to see details about the actual nameservers.

nameserver 192.168.101.254
nameserver 127.0.0.53
search lrde.epita.fr int.ionis-it.com
```

-  **8.8.8.8** is a popular public DNS server administrated by google, handy for testing purposes, do not use in production! 
- Mostly the same scheme on Windows systems
- Default server when using `dig`, `host`, ...

---

<sup>1</sup>it has to be an IP address and not a domain name.

# Exceptions: hosts file

Some exceptions can be stored in `/etc/hosts` file.

```
dstan@flan:~$ cat /etc/hosts
127.0.0.1 localhost
::1      ip6-localhost ip6-loopback
127.0.1.1 flan
```

Good practice: my local hostname (here: "flan") hardcoded as **127.0.0.1**.  
This is not **DNS queries** anymore:

```
dstan@flan:~$ host flan
Host flan not found: 3(NXDOMAIN)
dstan@flan:~$ getent hosts flan
127.0.1.1      flan
dstan@flan:~$ getent hosts epita.fr
2606:4700:20::ac43:4796 epita.fr
2606:4700:20::681a:7e1 epita.fr
```

Programmatic functions: `gethostbyname()`, `getaddrinfo()`<sup>2</sup>

---

<sup>2</sup><https://linux.die.net/man/3/gethostbyname>

# DNS Types

A DNS name can carry **multiple records**, of different **types**:

- A : IPv4 address
- AAAA : IPv6 address<sup>3</sup>
- NS : Name Server
- CNAME : Canonical Name (Alias)
- MX : Mail Exchanger, ie what mail server to use for this @domain
- ...

The following example has 2 A records, one AAAA, and one MX:

```
dstan@flan:~$ host epita.fr
epita.fr has address 104.26.6.225
epita.fr has address 104.26.7.225
epita.fr has IPv6 address 2606:4700:20::ac43:4796
epita.fr mail is handled by 0 epita-fr.mail.protection.outlook.com
```

---

<sup>3</sup>IPv4=32b=A hence IPv6=32 × 4 = 128b=AAAA

## Example: Query by type

Try it for yourself:

```
host -t CNAME www.lre.epita.fr
```

```
dig -t CNAME www.lre.epita.fr
```

```
host -t TXT lre.epita.fr
```

```
dig -t TXT lre.epita.fr
```

```
dig -t MX outlook.com
```

# Name Server type

```
dig -t NS ml.lre.epita.fr
# Nothing but:
dig -t NS lre.epita.fr
# lre.epita.fr name server ns.lre.epita.fr.
```

A Name Server (NS) server for *lre.epita.fr* is a DNS server **specifically** answering for the domain, but **nothing else**:

```
host intra.lre.epita.fr ns.lre.epita.fr
# intra.lre.epita.fr is an alias for rp.lre.epita.fr.
# rp.lre.epita.fr has address 91.243.117.236
# but:
host ionis-group.com ns.lre.epita.fr
# Host ionis-group.com not found: 5(REFUSED)
```

# Authoritative vs “Know-All”

- “*Know-All*” servers:
  - ▶ No matter what you ask them, they will answer.
  - ▶ Use them in your `/etc/resolv.conf`.
  - ▶ They’re anonymous: their IP won’t appear in NS answer query.
  - ▶ We will call them: **Resolver** or **Recursive Servers**
- **Authoritative** for a domain:
  - ▶ Everyone **knows**<sup>4</sup> they are authoritative for **this** domain (public knowledge).
  - ▶ If you ask them anything for another domain, they’ll **deny** you.

How to implement a recursive server?

---

<sup>4</sup>the domain name has a record of type NS to this authoritative server

## Recursive Resolution

# Recursive Resolution: An Example

Let's resolve *fleet.pie.cri.epita.fr* by querying the DNS server **198.97.190.53**

```
dig -t A fleet.pie.cri.epita.fr @198.97.190.53

# No answer for what I'm looking for, but ...
# .fr is managed by d.nic.fr
# We also have an additional section:
# (ADDITIONAL SECTION) d.nic.fr IN A 194.0.9.1

dig -t A fleet.pie.cri.epita.fr @194.0.9.1

# Still no answer, but now:
#          —> epita.fr has NS donna.ns.cloudflare.com:

dig -t A fleet.pie.cri.epita.fr @donna.ns.cloudflare.com

# Still no answer, but now:
#          —> cri.epita.fr ah NS ns.cri.epita.fr

dig -t A fleet.pie.cri.epita.fr @ns.cri.epita.fr


# An answer
# fleet.pie.cri.epita.fr is:
#   a CNAME to ingress.prod-1.k8s.cri.epita.fr.
#   which has IP 91.243.117.180
```

NB: NS entries answers have to be translated to IP addresses themselves!

- Initial server **198.97.190.53** was nice to us (ADDITIONAL SECTION) and gave us the IP of the NS: **glue record**.
- *d.nic.fr* was not so nice, we had to do a **second recursion** to resolve *donna.ns.cloudflare.com* to an IP (**173.245.58.151**).
- *donna.ns.cloudflare* was not so nice, we had to resolve again *ns.cri.epita.fr* to an IP (**91.243.117.210**).




# Recursive Resolution in a Nutshell



# Where do we start ? Root servers

For example: **198.97.190.53**

- Root servers IP addresses are **hardcoded**, publicly known by everyone.
- Full list: [https://en.wikipedia.org/wiki/Root\\_name\\_server#Root\\_server\\_addresses](https://en.wikipedia.org/wiki/Root_name_server#Root_server_addresses)
- They are **Widespread**: Several continents, AS, Companies
- In reality, **multiple** servers behind one IP
- Crucial for the Internet stability.



From arstechnica.com

<https://arstechnica.com/security/2024/05/dns-glitch-that-threatened-internet-stability-fixed-cause-remains-unclear/>

# Exercise: Recursive Resolution

Start with *j.root-servers.net* and give the NS servers used in the resolution for the below FQDNs.

NB: more than one solution.

- *www.epita.fr*
- *mail.dgfip.finances.gouv.fr*
- *ftp.debian.org*
- ...

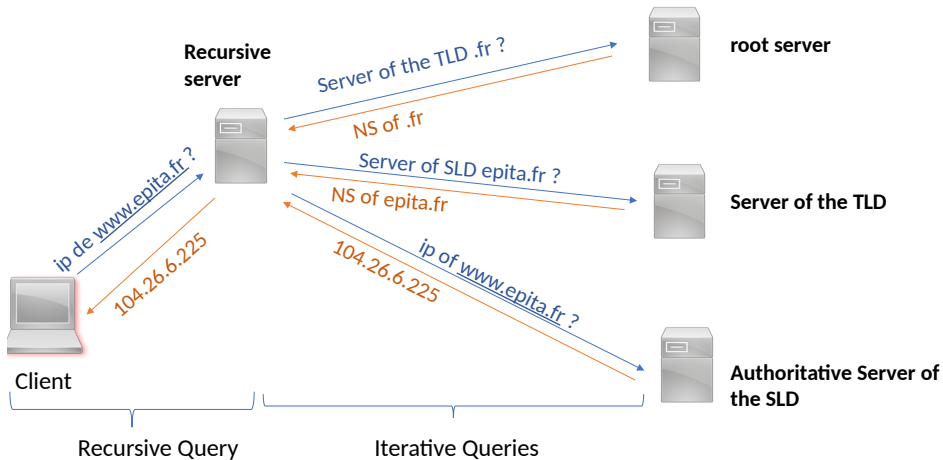
# Resilience and Caching

Some domains may have more than one NS, some may crash. (ex: *dns2.finances.gouv.fr.* at time of writing this slide). It is the admin's job to **synchronize all** their servers so they provide the same answer: usually one **main** authoritative and several **secondary** authoritative servers copying main.

- Record of type SOA (Start of Authority for a zone): gives the delay before any refresh
- Every DNS answer comes with a **TTL** (time to live).
- Based on this TTL, recursive servers **cache** results.
- **Relieves** the NS servers, especially root servers.

↪ editing DNS entries (to change IP) **needs to be planned** to avoid downtime.

# The global picture



## Remark: Recursive Server and Legal Implications

- DNS-blocking is the first technical measure to block a website in France, after a court order.
- ~→ operating a resolver in France may have **legal consequences**, especially if you are an ISP.
- Some reference: <https://www.bortzmeyer.org/censure-francaise.html>
- Another (telegram accidental block in May 2023):  
<https://www.bortzmeyer.org/blocage-telegram-france.html>

```
# NB: 192.168.1.254 is my ISP modem
dstan@flan:~$ host www.palaceofchance.com 192.168.1.254

www.???ofchance.com is an alias for offre-illegale.anj.fr.
offre-illegale.anj.fr has address 145.239.225.117

dstan@flan:~$ host www.???ofchance.com 8.8.8.8

www.???ofchance.com has address 137.????.????.???
```

# Summary

- The DNS is a giant global **database** storing key/values;
- Values can be of **many types** (A,AAAA,NS, ...)
- The keys are FQDNs, which are **hierachically delegated**: sub<SLD<TLD<"."
- Two kinds of DNS servers: **authoritative** and **recursive**
- Authoritative: answers for **one domain only**
- Exercise: perform **recursive resolution** by hand, from root to the final NS.
- Conclusion: **Distributed system, Vital** for the Internet.