

# Basics of Networking 2: Firewalls and NAT

Daniel STAN




June 19, 2024

# Firewall: Motivation

Main motivation: increase security

- Prevent **unauthorized connections**
- Rate **limit communications** (against DOS)
- Avoid **exploitation of vulnerabilities**.



NB: these do not prevent vulnerabilities in software, just **reduce the exposure**.

# Goals for today

- (Briefly) analyze a machine/network surface of attack
- Write (small) firewall rules

We will illustrate all these concepts with the **linux firewall** system. More precisely: **iptables**.

# Goals for today

- (Briefly) analyze a machine/network surface of attack
- Write (small) firewall rules

We will illustrate all these concepts with the **linux firewall** system. More precisely: **iptables**.

We will study connection tracking, as a **by-product** of the firewall system:

- How to track TCP sessions?
- How to use it to perform **Network Address Translation**?

# You are at risk

List communications on a Linux system: `netstat -ltuapen`  
 (l=listen, u=udp, t=tcp, n=no-dns, p=processes)

```
$ sudo netstat -ltuapen
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address           Foreign Address         State       PID/Program name
tcp        0      0 127.0.0.1:9050          0.0.0.0:*                LISTEN      1807/tor
tcp        0      0 0.0.0.0:57621          0.0.0.0:*                LISTEN      1806861/spotify
tcp        0      0 192.168.170.217:45690  34.128.128.0:443        ESTABLISHED 835950/chrome —typ
# ... 67 lines
```

Are all these communications legit, or should some of them be blocked?

# Wireshark

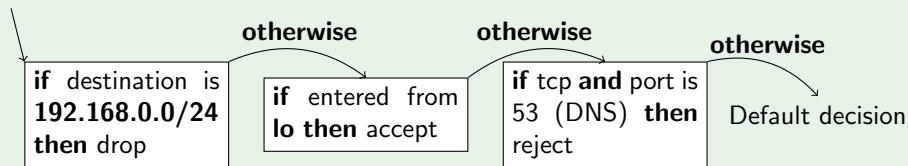
As opposed to `netstat`, `wireshark` also references communication not for our own machine, but also packets routed for others.

# IPTables

iptables is a firewall linux system composed of

- Several **rules**, stored in a chain.
- Several **chains**, basically just a linked list, chains are stored in a table.
- Several **tables**: *filter*, *nat*, *mangle*.

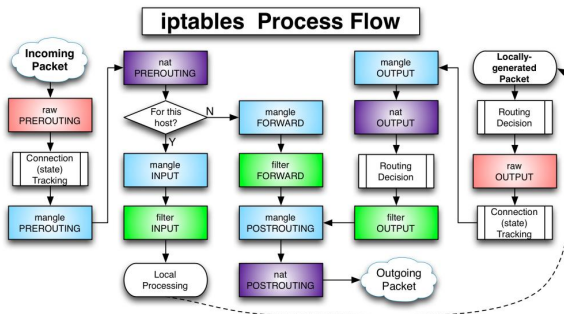
## Example



Informal example of a chain made of 3 rules.

Each IP packet is processed by a chain, when some specific **network event** occurs.

# Network Events: Flow of a packet in iptables



Ref: <https://westoahu.hawaii.edu/cyber/wp-content/uploads/2017/11/iptables-Flowchart.jpg>

- Locally processed = a process on the computer.
- Chains are written in UPPER CASE, tables in small case.
- We mostly use `filter` (default table).



# The filter table

This is the default table, it contains three chains:

- INPUT: incoming packets
- OUTPUT: outgoing packet
- FORWARD: packets which are only passing by

# The filter table

This is the default table, it contains three chains:

- INPUT: incoming packets
- OUTPUT: outgoing packet
- FORWARD: packets which are only passing by



**Applicative** firewall



**Router's** firewall

# Decision: Targets

Taking a decision for a packet means to send the packet to another *special* chain:

- ACCEPT accepts the packet, without further reading the chain rules;
- DROP drops the packet, **silently**;
- REJECT same as DROP but also sends a message back (ICMP error) to the sender.

**Rule of thumb:** if you don't want to disclose your internal network, just drop packets, but rejecting is better for debugging purposes.

# Decision: Targets

Taking a decision for a packet means to send the packet to another *special* chain:

- ACCEPT accepts the packet, without further reading the chain rules;
- DROP drops the packet, **silently**;
- REJECT same as DROP but also sends a message back (ICMP error) to the sender.

**Rule of thumb:** if you don't want to disclose your internal network, just drop packets, but rejecting is better for debugging purposes.

For example, to reject all routed packets through our machine:

```
# -A = appends (queue of the list), -j = jump  
iptables -A FORWARD -j REJECT
```


# Conditions: Interfaces

Plenty of conditions: <https://linux.die.net/man/8/iptables>

## Network Interfaces

`-i` (input) followed by an interface name specifies the interface used by the packet to enter the system, while `-o` (output) specifies the planned interface for outputting the packet.

```
# accept all packets routed from eth0 AND to eth1
iptables -A FORWARD -i eth0 -o eth1 -j ACCEPT
# reject all the rest
iptables -A FORWARD -j REJECT
```



# Conditions: Interfaces

Plenty of conditions: <https://linux.die.net/man/8/iptables>

## Conditions: IP Address or Subnet

--src specifies the source IP address (or subnet) while --dst specifies the destination address (or subnet).

```
# throws away packets
# *from eth0 AND
# *with source IP in 192.168.0.0/16
iptables -A FORWARD --src 192.168.0.0/24 \
        -i eth0 -j DROP
```

This takes the conjunction (AND). One can also take negations with ! keyword.

# Conditions: Protocol

## Transport Protocol

`-p` specifies a transport protocol (above IP), that is to say for example `tcp`, `udp` or `icmp` (ping). (one could also use the protocol number instead of its name).  
If `tcp` or `udp`, one can specify ports with `--dport` or `--sport`.

```
# Let IP packets of type TCP pass by
iptables -I FORWARD -p tcp -j ACCEPT
# Authorize a SSH (22/tcp) server on this computer
iptables -I INPUT -p tcp --dport 22 -j ACCEPT
```

## (optional) Custom Chains

Chains are some kind of functions in a programming language, hence the `RETURN` target. Example of a custom chain that be re-used:

```
# -N = new chain
-N SUSPICIOUS
-A SUSPICIOUS -p udp --dport 53 -j DROP
-A SUSPICIOUS -p tcp --dport 22 -j DROP
-A SUSPICIOUS -p tcp --dport 80 -j DROP
-A SUSPICIOUS -j RETURN

-A FORWARD -i eth0 -j SUSPICIOUS
-A FORWARD -i eth1 -j SUSPICIOUS
-A FORWARD -j ACCEPT
```



## (optional) Custom Chains

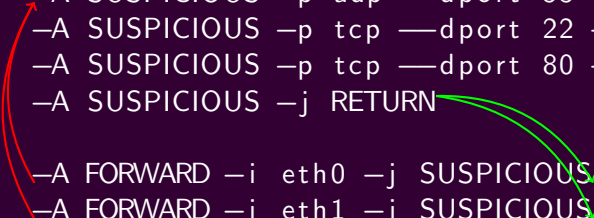
Chains are some kind of functions in a programming language, hence the `RETURN` target. Example of a custom chain that be re-used:

```
# -N = new chain
-N SUSPICIOUS
-A SUSPICIOUS -p udp --dport 53 -j DROP
-A SUSPICIOUS -p tcp --dport 22 -j DROP
-A SUSPICIOUS -p tcp --dport 80 -j DROP
-A SUSPICIOUS -j RETURN
-A FORWARD -i eth0 -j SUSPICIOUS
-A FORWARD -i eth1 -j SUSPICIOUS
-A FORWARD -j ACCEPT
```

## (optional) Custom Chains

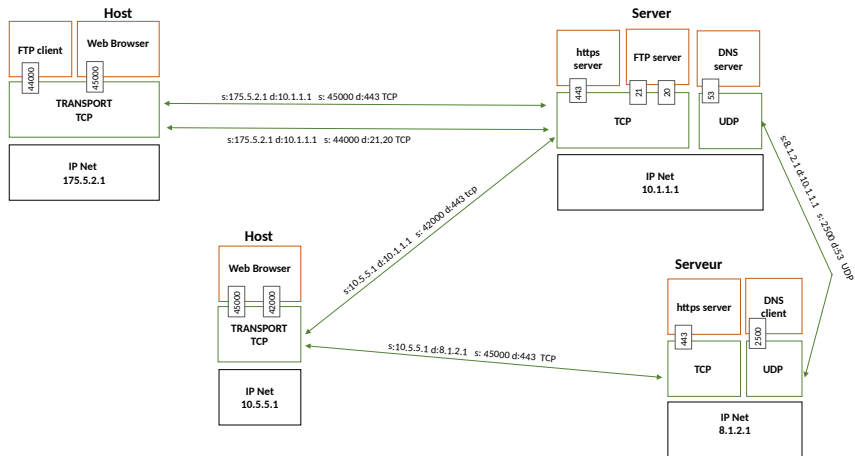
Chains are some kind of functions in a programming language, hence the `RETURN` target.  
Example of a custom chain that be re-used:

```
# -N = new chain
-N SUSPICIOUS
-A SUSPICIOUS -p udp --dport 53 -j DROP
-A SUSPICIOUS -p tcp --dport 22 -j DROP
-A SUSPICIOUS -p tcp --dport 80 -j DROP
-A SUSPICIOUS -j RETURN
-A FORWARD -i eth0 -j SUSPICIOUS
-A FORWARD -i eth1 -j SUSPICIOUS
-A FORWARD -j ACCEPT
```



## Stateful Firewalls

# Recall: TCP sessions, made of several IP packets



# Stateful vs Stateless

So far, we inspected IP packets individually, but they may be part of a TCP session, so it might be best to filter them together. This requires the firewall to have some **memory** of what previous packets passed by:

- Stateless: filters packets **individually**;
- Stateful: some **memory**.

# States with iptables

## Condition: TCP/UDP state

With `-m state`, one can specify the state of the by `--state` and a set of possible connection states.

Possible state types: NEW, ESTABLISHED, RELATED, INVALID, UNTRACKED.

```
iptables -A FORWARD -m state \  
    --state RELATED,ESTABLISHED -j ACCEPT  
  
# This second expensive rule will be inspected  
# only for new connections, and not for every packet  
iptables -A FORWARD -m state --state NEW \  
    -j EXPENSIVE_ANALYSIS
```

# How a firewall records a TCP session?

On linux, `conntrack -L` displays the memory of the firewall, that is to say the connection tracking mechanism:

```
$ conntrack -L  
#...
```

A connection is **uniquely identified** by 4 informations: (Source, Destination)  $\times$  (IP, Port).

**Recall:** when connecting to a server, the destination port is known (80 for http, 443 for https, etc), but the source port is **chosen at random** by the OS.

# Network Address Translation (NAT)


The quadruple (Source, Destination)  $\times$  (IP, Port) is stored **twice** in conntrack: once for the **input** network interface and another for the **output** interface.

→ **NAT**: The firewall can *rewrite* IP addresses/ports, in a consistent way.



# NAT: example and usage


- Private addresses: **192.168.42.0/24**, cannot be used on the internet as they are **not unique**.
- Only a single public IP **82.235.136.108**.



Say M2 connects to <http://104.26.7.225>,

# NAT: example and usage

- Private addresses: **192.168.42.0/24**, cannot be used on the internet as they are **not unique**.
- Only a single public IP **82.235.136.108**.




Say M2 connects to `http://104.26.7.225`,

- Query packet
  - ▶ On the local network: (**192.168.42.22**,**104.26.7.225**,15077,80)

# NAT: example and usage

- Private addresses: **192.168.42.0/24**, cannot be used on the internet as they are **not unique**.
- Only a single public IP **82.235.136.108**.




Say M2 connects to `http://104.26.7.225`,

- Query packet
  - ▶ On the local network: (**192.168.42.22**,**104.26.7.225**,15077,80)
  - ▶ On the internet, after NAT: (**82.235.136.108**,**104.26.7.225**,15077,80)

# NAT: example and usage

- Private addresses: **192.168.42.0/24**, cannot be used on the internet as they are **not unique**.
- Only a single public IP **82.235.136.108**.




Say M2 connects to `http://104.26.7.225`,

- Answer from the HTTP server:
  - ▶ On the internet: (**104.26.7.225,82.235.136.108,80,15077**)

# NAT: example and usage


- Private addresses: **192.168.42.0/24**, cannot be used on the internet as they are **not unique**.
- Only a single public IP **82.235.136.108**.



Say M2 connects to `http://104.26.7.225`,

- Answer from the HTTP server:
  - ▶ On the internet: (**104.26.7.225**, **82.235.136.108**, 80, 15077)
  - ▶ On the local network, after NAT: (**104.26.7.225**, **192.168.42.22**, 80, 15077)


# NAT: example and usage



Say M2 and M3 both connect to `http://104.26.7.225`:

- Query packets
  - ▶ On the local network (M2): (**192.168.42.22,104.26.7.225,929,80**)
  - ▶ (M3): (**192.168.42.33,104.26.7.225,7207,80**)

# NAT: example and usage



Say M2 and M3 both connect to `http://104.26.7.225`:

- Query packets

- ▶ On the local network (M2): (192.168.42.22, 104.26.7.225, 929, 80)
- ▶ (M3): (192.168.42.33, 104.26.7.225, 7207, 80)
- ▶ On the internet, after NAT: (82.235.136.108, 104.26.7.225, 929, 80)
- ▶ (82.235.136.108, 104.26.7.225, 7207, 80)

# NAT: example and usage

Say M2 and M3 both connect to `http://104.26.7.225`:

- Query packets
  - ▶ On the local network (M2): (**192.168.42.22**,**104.26.7.225**,929,80)
  - ▶ (M3): (**192.168.42.33**,**104.26.7.225**,7207,80)
  - ▶ On the internet, after NAT: (**82.235.136.108**,**104.26.7.225**,929,80)
  - ▶ (**82.235.136.108**,**104.26.7.225**,7207,80)
- Answers from the HTTP server:
  - ▶ On the internet: (**104.26.7.225**,**82.235.136.108**,80,929)
  - ▶ (**104.26.7.225**,**82.235.136.108**,80,7207)



# NAT: example and usage

Say M2 and M3 both connect to `http://104.26.7.225`:

- Query packets
  - ▶ On the local network (M2): (**192.168.42.22**,**104.26.7.225**,929,80)
  - ▶ (M3): (**192.168.42.33**,**104.26.7.225**,7207,80)
  - ▶ On the internet, after NAT: (**82.235.136.108**,**104.26.7.225**,929,80)
  - ▶ (**82.235.136.108**,**104.26.7.225**,7207,80)
- Answers from the HTTP server:
  - ▶ On the internet: (**104.26.7.225**,**82.235.136.108**,80,929)
  - ▶ (**104.26.7.225**,**82.235.136.108**,80,7207)
  - ▶ On the local net, after NAT: (**104.26.7.225**,**192.168.42.22**,80,929)
  - ▶ (**104.26.7.225**,**192.168.42.33**,80,7207)

# NAT: example and usage

Say M2 and M3 both connect to `http://104.26.7.225`:

- Query packets
  - ▶ On the local network (M2): (192.168.42.22,104.26.7.225,929,80)
  - ▶ (M3): (192.168.42.33,104.26.7.225,7207,80)
  - ▶ On the internet, after NAT: (82.235.136.108,104.26.7.225,929,80)
  - ▶ (82.235.136.108,104.26.7.225,7207,80)
- Answers from the HTTP server:
  - ▶ On the internet: (104.26.7.225,82.235.136.108,80,929)
  - ▶ (104.26.7.225,82.235.136.108,80,7207)
  - ▶ On the local net, after NAT: (104.26.7.225,192.168.42.22,80,929)
  - ▶ (104.26.7.225,192.168.42.33,80,7207)

**Q:** What if the same port is picked at random (*bad luck*)?

# NAT: example and usage

Say M2 and M3 both connect to `http://104.26.7.225`:

- Query packets
  - ▶ On the local network (M2): (192.168.42.22,104.26.7.225,929,80)
  - ▶ (M3): (192.168.42.33,104.26.7.225,7207,80)
  - ▶ On the internet, after NAT: (82.235.136.108,104.26.7.225,929,80)
  - ▶ (82.235.136.108,104.26.7.225,7207,80)
- Answers from the HTTP server:
  - ▶ On the internet: (104.26.7.225,82.235.136.108,80,929)
  - ▶ (104.26.7.225,82.235.136.108,80,7207)
  - ▶ On the local net, after NAT: (104.26.7.225,192.168.42.22,80,929)
  - ▶ (104.26.7.225,192.168.42.33,80,7207)

Q: What if the same port is picked at random (*bad luck*)?

→ NAT can also **rewrite ports**

# Linux: the nat table

Option `-t nat`, this table has four chains:

- PREROUTING, INPUT
- POSTROUTING, OUTPUT

# Linux: the nat table

Option `-t nat`, this table has four chains:

- PREROUTING, INPUT
- POSTROUTING, OUTPUT

Three new targets:

- SNAT: rewrites the source IP together with `--to-source`
- DNAT: rewrites the destination IP of a session, together with `--to-destination`

Exercise/Question: what rule was used on the previous example?

# Linux: the nat table

Option `-t nat`, this table has four chains:

- PREROUTING, INPUT
- POSTROUTING, OUTPUT

Three new targets:

- SNAT: rewrites the source IP together with `--to-source`
- DNAT: rewrites the destination IP of a session, together with `--to-destination`

Exercise/Question: what rule was used on the previous example?

```
iptables -t nat -A POSTROUTING --src 192.168.42.0/24 \  
-j SNAT --to-source 82.235.136.109
```

## (optional) MASQUERADE target

MASQUERADE is a shortcut for SNAT where the source address is **automatically computed** using the routes.

### Example

Assume *R1* has three IPs **192.168.42.1**, **82.235.136.108** and **10.0.0.1**.

It has the following routes:

```
default via 82.235.136.1 dev eth0 src 82.235.136.108
10.0.0.0/24 dev eth1 src 10.0.0.1
192.168.42.0/24 dev eth2 src 192.168.42.1
```

Then `-j MASQUERADE` is equivalent to:

```
-A POSTROUTING -o eth0 -j SNAT --to-source 82.235.136.108
```

## (optional) MASQUERADE target

MASQUERADE is a shortcut for SNAT where the source address is **automatically computed** using the routes.

### Example

Assume *R1* has three IPs **192.168.42.1**, **82.235.136.108** and **10.0.0.1**.

It has the following routes:

```
default via 82.235.136.1 dev eth0 src 82.235.136.108
10.0.0.0/24 dev eth1 src 10.0.0.1
192.168.42.0/24 dev eth2 src 192.168.42.1
```

Then `-j MASQUERADE` is equivalent to:

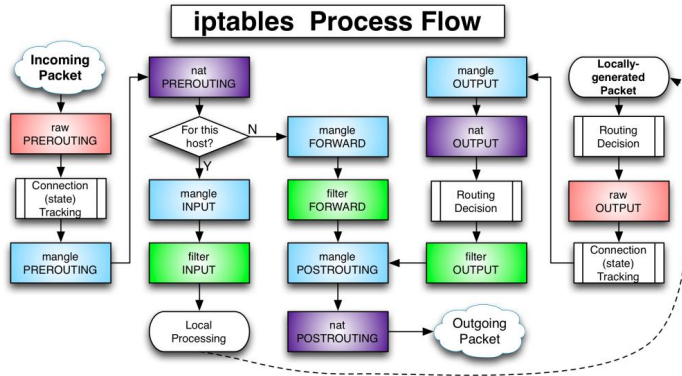
```
-A POSTROUTING -o eth0 -j SNAT --to-source 82.235.136.108
-A POSTROUTING -o eth1 -j SNAT --to-source 10.0.0.1
-A POSTROUTING -o eth2 -j SNAT --to-source 192.168.42.1
```



# Recall: Network Flow

- SNAT is to be used **after** routing decision: valid only in POSTROUTING and OUTPUT;
- DNAT is to be used **before** routing decision: valid only in PREROUTING and INPUT.

NB: this also means that `-i` cannot be used for SNAT and `-o` cannot be used for DNAT.



# DNAT: MITM or caching

DNAT: used to intercept communications. For example, one could intercept some HTTP connection to serve it with a local server:

```
# Let's make 192.168.42.2 believe he's talking to ftp.debian.org:80
# while in fact, he's talking to me
-A PREROUTING —src 192.168.42.2 —dst 199.232.170.132 \
    -p tcp —dport 80 \
    -j DNAT —to-destination 10.0.0.1:80
```

Useful technique to introduce a **local mirrors**, **captive portals** on a Wifi hotspot, or to perform **Man-in-the-middle** attacks.

NB: doesn't work well with **authenticated communications**, for example https.

# Exercise: NAT and conntrack

Analyze the following conntrack -l output, and give the different used NAT rules:

dev1	src1	dst1	dev2	src2	dst2
<b>eth0</b>	192.168.0.3:7331	146.75.118.132:80	<b>eth2</b>	82.235.136.108:7331	146.75.118.132:80
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:2222	<b>eth0</b>	172.67.71.150:4242	192.168.0.2:22
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:3333	<b>eth0</b>	172.67.71.150:4242	192.168.0.3:22

# Exercise: NAT and conntrack

Analyze the following conntrack -l output, and give the different used NAT rules:

dev1	src1	dst1	dev2	src2	dst2
<b>eth0</b>	192.168.0.3:7331	146.75.118.132:80	<b>eth2</b>	82.235.136.108:7331	146.75.118.132:80
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:2222	<b>eth0</b>	172.67.71.150:4242	192.168.0.2:22
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:3333	<b>eth0</b>	172.67.71.150:4242	192.168.0.3:22

- SNAT

# Exercise: NAT and conntrack

Analyze the following conntrack -l output, and give the different used NAT rules:

dev1	src1	dst1	dev2	src2	dst2
<b>eth0</b>	192.168.0.3:7331	146.75.118.132:80	<b>eth2</b>	82.235.136.108:7331	146.75.118.132:80
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:2222	<b>eth0</b>	172.67.71.150:4242	192.168.0.2:22
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:3333	<b>eth0</b>	172.67.71.150:4242	192.168.0.3:22

- SNAT  
--src 192.168.0.3 -j SNAT \$PUB
- DNAT

# Exercise: NAT and conntrack

Analyze the following conntrack -l output, and give the different used NAT rules:

dev1	src1	dst1	dev2	src2	dst2
<b>eth0</b>	192.168.0.3:7331	146.75.118.132:80	<b>eth2</b>	82.235.136.108:7331	146.75.118.132:80
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:2222	<b>eth0</b>	172.67.71.150:4242	192.168.0.2:22
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:3333	<b>eth0</b>	172.67.71.150:4242	192.168.0.3:22

- SNAT

```
--src 192.168.0.3 -j SNAT $PUB
```

- DNAT

```
--dst $PUB -p tcp --dport 2222 -j DNAT --to-destination 192.168.0.2:22
```

# Exercise: NAT and conntrack

Analyze the following conntrack -l output, and give the different used NAT rules:

dev1	src1	dst1	dev2	src2	dst2
<b>eth0</b>	192.168.0.3:7331	146.75.118.132:80	<b>eth2</b>	82.235.136.108:7331	146.75.118.132:80
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:2222	<b>eth0</b>	172.67.71.150:4242	192.168.0.2:22
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:3333	<b>eth0</b>	172.67.71.150:4242	192.168.0.3:22

- SNAT  
 --src 192.168.0.3 -j SNAT \$PUB
- DNAT  
 --dst \$PUB -p tcp --dport 2222 -j DNAT --to-destination 192.168.0.2:22
- DNAT

# Exercise: NAT and conntrack

Analyze the following conntrack -l output, and give the different used NAT rules:

dev1	src1	dst1	dev2	src2	dst2
<b>eth0</b>	192.168.0.3:7331	146.75.118.132:80	<b>eth2</b>	82.235.136.108:7331	146.75.118.132:80
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:2222	<b>eth0</b>	172.67.71.150:4242	192.168.0.2:22
<b>eth2</b>	172.67.71.150:4242	82.235.136.108:3333	<b>eth0</b>	172.67.71.150:4242	192.168.0.3:22

- SNAT  
--src 192.168.0.3 -j SNAT \$PUB
- DNAT  
--dst \$PUB -p tcp --dport 2222 -j DNAT --to-destination 192.168.0.2:22
- DNAT  
--dst \$PUB -p tcp --dport 3333 -j DNAT --to-destination 192.168.0.3:22

NB: typical setup with multiple SSH servers behind NAT, by using different ports on the public IP.



# NAT Summary

- SNAT: heavily used to overcome the **IP shortage problem**:
  - ▶ Every customer gets only **a single** public IP for all its devices, the router (box) performs SNAT
  - ▶ Carrier-Grade NAT (CGNAT): same but at the ISP level.
- DNAT:
  - ▶ MITM attacks;
  - ▶ Captive Portals;
  - ▶ Host servers behind your ISP modem;

Recall: NAT requires tracking all connections, this can become **very expensive** in a large network.

# Epilogue

In this class and in lab sessions, you have seen:

- How to setup a local network with `ip a`, `ip r`
- Distribute IP addresses (DHCP) and provide DNS services with `dnsmasq`
- How to share a single IP address with `iptables`

This is exactly how most ISP modem work, on GNU/Linux, for example

`http://floss.freebox.fr:`

Divers / Mentions légales			
La Freebox Server utilise des logiciels libres ou opensource. La liste est présente ci-dessous. Vous pouvez obtenir le code source complet des logiciels concernés à l'adresse suivante: <a href="http://floss.freebox.fr/">http://floss.freebox.fr/</a>			
Nom	Version	Licence	Url
dnsmasq	2.85	GPLv2	<a href="http://floss.freebox.fr/packages/dnsmasq-2.85.tar.gz">http://floss.freebox.fr/packages/dnsmasq-2.85.tar.gz</a>
dosfstools	2.11	GPLv2	<a href="http://floss.freebox.fr/packages/dosfstools-2.11.src.tar.gz">http://floss.freebox.fr/packages/dosfstools-2.11.src.tar.gz</a>
dropbear	2022.82	MIT	
e2fsprogs	1.42.9	GPLv2	<a href="http://floss.freebox.fr/packages/e2fsprogs-1.42.9.tar.gz">http://floss.freebox.fr/packages/e2fsprogs-1.42.9.tar.gz</a>
ethtool	5.3	GPLv2	<a href="http://floss.freebox.fr/packages/ethtool-5.3.tar.gz">http://floss.freebox.fr/packages/ethtool-5.3.tar.gz</a>
expat	2.4.8	BSD	
ffmpeg	4.3.1	LGPL	<a href="http://floss.freebox.fr/packages/ffmpeg-4.3.1.tar.xz">http://floss.freebox.fr/packages/ffmpeg-4.3.1.tar.xz</a>